

Bayesian Hierarchical Modeling of COVID-19 Cases and Government Response in the United States

Tianshu Liu, Jiong Ma, Jiayi Shi, Zhengwei Song, Ziqing Wang

April 2023

Contents

1	Background	1
2	Data Processing	1
3	Method	2
3.1	Bayesian Hierarchical Model	2
3.1.1	Hierarchical Structure	3
3.1.2	Choices of Priors	3
3.1.3	Likelihood	4
3.1.4	Posterior	4
3.2	Component-wise MH Algorithm	5
3.3	MCMC Chain Convergence Diagnostics	6
3.3.1	Diagnostic Plots	6
3.3.2	Gelman-Rubin Statistic	6
4	Result	7
4.1	Convergence of MCMC Chain	7
4.2	Posterior Distribution	13
4.3	Credible Intervals	15
5	Discussion	15
5.1	Interpretation	15
5.2	Public Health Policy Recommendations	17
5.3	Suggestions on Further Research	17
6	Conclusion	17
Acknowledgement		17
Reference		17
Code Implementation of MCMC		18

1 Background

The COVID-19 pandemic has had a profound impact on the world, causing significant disruptions in public health, economic stability, and social life. In the United States, the pandemic has highlighted the importance of understanding the spread of the virus and the effectiveness of various government responses in mitigating its impact. As the pandemic unfolded, different states implemented a wide range of containment measures, economic support policies, and public health interventions. Understanding the relationship between these factors and the spread of COVID-19 is crucial for informing future policy decisions and improving public health outcomes.

To investigate this relationship, this project aims to apply Bayesian hierarchical modeling techniques to analyze the COVID-19 case data and various government response indexes in the United States during the year 2020 before the vaccine became available. The government response indexes include containment measures, economic support policies, and stringency levels. It also investigates the role of demographic factors, such as population density and elderly population, in the spread of COVID-19 and the effectiveness of government responses.

In addition, this project also aims to obtain a good understanding of the factors influencing the pandemic's impact on public health. Ultimately, the insights gained from this analysis can help inform future policy decisions and guide effective response strategies in the face of ongoing and future public health emergencies.

2 Data Processing

The original dataset is created by merging multiple sources of information to facilitate the Bayesian hierarchical modeling of COVID-19 cases and government responses in the United States. It comprises weekly aggregated data for the year 2020, including state-level COVID-19 case counts, government response indexes (Government Response Index, Containment Health Index, Economic Support Index, and Stringency Index), and mobility changes in retail, parks, and transit stations. Additionally, the dataset incorporates state-level demographic information, such as population density and the proportion of the elderly population. The data provides a good picture of the pandemic landscape in the United States in 2020.

The original dataset has 2548 observations and 15 variables: state, week start, weekly cases, pop2019, LandArea, Percentage over 65, infection rate, retail and recreation percent change from baseline, parks percent change from baseline, transit stations percent change from baseline, population density, government response index, containment index, economic support index, stringency index. Table 1 gives the summary statistics before scaling.

Table 1: Summary statistics before scaling

Statistic	N	Mean	St. Dev.	Min	Max
retail_and_recreation_percent_change_from_baseline	2,327	-10.187	14.819	-59.877	23.707
parks_percent_change_from_baseline	2,327	46.095	61.362	-70.286	407.000
transit_stations_percent_change_from_baseline	2,327	-12.803	18.451	-77.714	55.250
population_density	2,327	206.017	270.013	1.280	1,254.244
government_response_index	2,327	54.757	16.550	9.006	80.210
containment_index	2,327	55.242	15.770	10.290	79.640
economic_support_index	2,327	51.358	27.870	0.000	100.000
stringency_index	2,327	57.465	18.656	6.744	93.520

The data processing step is then done in 3 directions.

1. Variable inspection: The weekly state-level cumulative case counts is transformed to new case counts. This is to correspond with the assumption of further Possion model.
2. Variable selection: Only four variables were selected as covariates for analysis, which include the Government Response Index, weekly average percentage change in mobility trends for Retail and Recreation Places, Parks, and Transit Stations in each state. This decision is made based on the high correlation observed among several variables in Figure 1, including the Government Response Index, Containment Health Index, Economic Support Index, and Stringency Index. Including highly correlated variables can result in longer computation times and may cause Bayesian hierarchical models to fail to converge when estimating parameters. In addition, it is found that the Containment Health Index, Economic Support Index, and Stringency Index variables could be interpreted using the Government Response Index.

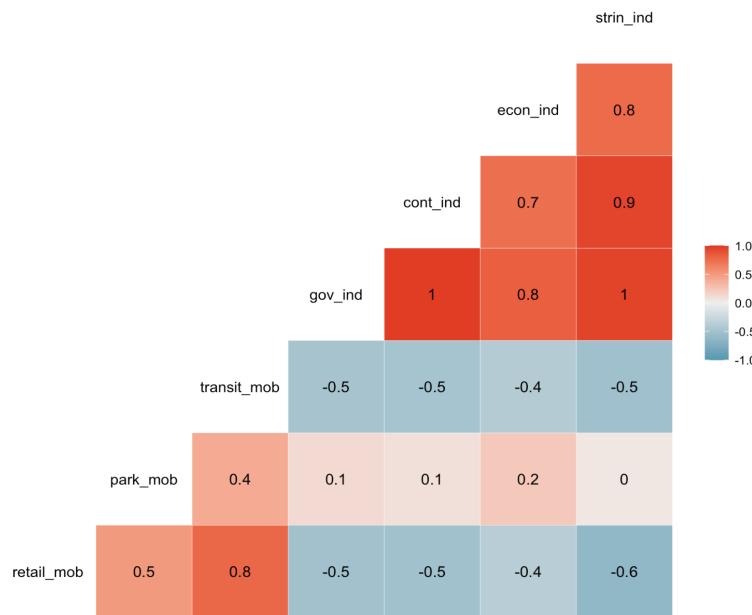


Figure 1: Pair-wise correlations between covariates in the data set

3. Variable scaling: The state-level COVID-19 case counts and Population in 2019 is measured in units of millions. To scale the other predictors, they are divided by their respective standard deviations. Scaling helps to ensure that all predictors have equal weight in the model, so they can be more fairly compared to one another. Also, scaling can help to improve the numerical stability of the model. Many numerical optimization algorithms used in Bayesian hierarchical models, such as Markov Chain Monte Carlo (MCMC), can be sensitive to the scale of the predictors.

3 Method

3.1 Bayesian Hierarchical Model

A Bayesian hierarchical model is a statistical model that allows for modeling complex data structures by incorporating multiple levels of variation. In this model, the data is assumed to be generated from a hierarchical structure, where each level of the hierarchy corresponds to a different level of variation in the data.

The parameters of the model are estimated by fitting the model to the observed data using Bayesian inference:

$$g(\lambda|Y) = \frac{f(Y|\lambda)\pi(\lambda)}{f(Y)}$$

This involves specifying a likelihood function $f(Y|\lambda)$ that describes the probability of observing the data given the model parameters, as well as prior distributions $\pi(\lambda)$ that

describe our prior knowledge about the model parameters. The posterior distribution of the parameters $g(\lambda|Y)$ is then obtained. This posterior distribution can be used to make predictions about future observations, as well as to estimate the uncertainty in the model parameters.

Bayesian hierarchical models allow for the incorporation of prior information into the analysis, can handle complex data structures with multiple levels of variation, and can provide more accurate estimates of uncertainty in the model parameters.

In this study, we assume that the number of new infections in state i during week j , denoted as Y_{ij} , follows a Poisson distribution:

$$Y_{ij} \sim \text{Poisson}(\lambda_{ij} n_{ij})$$

$$P(Y_{ij} = k) = \frac{e^{-\lambda_{ij} n_{ij}} (\lambda_{ij} n_{ij})^k}{k!}$$

where λ_{ij} is the infection rate in state i during week j , and n_{ij} is the population of state i during week j .

3.1.1 Hierarchical Structure

We model the log of infection rates $\log(\lambda_{ij})$ using a linear model that includes state-level random effects and fixed effects for the covariates:

$$\log(\lambda_{ij}) = \alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon$$

$$\lambda_{ij} = \exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon)$$

where α is the overall intercept, β is a vector of fixed effects coefficients for the covariates X_{ij} (e.g., government response index, and mobility changes), γ is the fixed effect coefficient for the population density P_{ij} , δ is the fixed effect coefficient for the percentage of the elderly population E_{ij} , u_i is the state-level random effect, and ϵ_{ij} is the residual error term.

3.1.2 Choices of Priors

We choose weakly informative priors for the fixed effects coefficients β_k , γ , and δ , such as normal distributions with mean 0 and a large variance:

$$\beta_k \sim \text{Normal}(0, 10^2)$$

$$\gamma \sim \text{Normal}(0, 10^2)$$

$$\delta \sim \text{Normal}(0, 10^2)$$

For the overall intercept α , we use a weakly informative prior:

$$\alpha \sim \text{Normal}(0, 10^2)$$

For the state-level random effects u_i , we assume that they follow a normal distribution with mean 0 and a common variance σ_u^2 :

$$u_i \sim \text{Normal}(0, \sigma_u^2)$$

We choose a weakly informative prior for the standard deviation σ_u :

$$\sigma_u \sim \text{HalfNormal}(0, 10^2) \quad \text{or} \quad \sigma_u \sim \text{HalfCauchy}(0, 10)$$

For the residual error term ϵ_{ij} , we assume that it follows a normal distribution with mean 0 and a common variance σ_ϵ^2 :

$$\epsilon_{ij} \sim \text{Normal}(0, \sigma_\epsilon^2)$$

We choose a weakly informative prior for the standard deviation σ_ϵ :

$$\sigma_\epsilon \sim \text{HalfNormal}(0, 10^2) \quad \text{or} \quad \sigma_\epsilon \sim \text{HalfCauchy}(0, 10)$$

Let the parameter space Θ to denote the set of parameters. Thus,

$$\Theta = \{\alpha, \beta_1, \dots, \beta_4, \gamma, \delta, u_1, \dots, u_{50}, \epsilon, \sigma_u, \sigma_\epsilon\}$$

Let n_s denote the total number of states, the prior can be calculated using:

$$\begin{aligned} \pi(\Theta) &= \pi(\alpha) \prod_{i=1}^4 \pi(\beta_i) \pi(\gamma) \pi(\delta) \prod_{i=1}^{n_s} \pi(u_i | \sigma_u) \pi(\sigma_u) \pi(\epsilon | \sigma_\epsilon) \pi(\epsilon) \\ &\propto \exp\left\{-\frac{\alpha^2 - \sum_{i=1}^k \beta_k^2 - \sigma_u^2 - \sigma_\epsilon^2}{2 \cdot 10^2}\right\} \frac{1}{\sigma_u \sigma_\epsilon} \exp\left\{-\frac{\sum_{i=1}^{n_s} u_i^2}{2\sigma_u^2}\right\} \exp\left\{-\frac{\epsilon^2}{2\sigma_\epsilon^2}\right\} \end{aligned}$$

$$\begin{aligned} \log \pi(\Theta) &= \log \pi(\alpha) + \sum_{i=1}^4 \log \pi(\beta_i) + \log \pi(\gamma) + \log \pi(\delta) + \sum_{i=1}^{n_s} \log \pi(u_i | \sigma_u) \\ &\quad + \log \pi(\sigma_u) + \log \pi(\epsilon | \sigma_\epsilon) + \log \pi(\epsilon) \\ &\propto \frac{-\alpha^2 - \sum_{i=1}^k \beta_k^2 - \sigma_u^2 - \sigma_\epsilon^2}{2 \cdot 10^2} - \log(\sigma_u) - \log(\sigma_\epsilon) - \frac{\sum_{i=1}^{n_s} u_i^2}{2\sigma_u^2} - \frac{\epsilon^2}{2\sigma_\epsilon^2} \end{aligned}$$

3.1.3 Likelihood

Let n_s denote the total number of states and n_w denote the total number of weeks. Since the number of new infections in state i during week j , denoted as y_{ij} , follows a Poisson distribution: $y_{ij} \sim \text{Poisson}(\lambda_{ij} n_{ij})$ where λ_{ij} is the infection rate in state i during week j , and n_{ij} is the population of state i during week j , the likelihood can be calculated as:

$$\begin{aligned} \lambda_{ij} &= \exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon) \\ L_Y(\Theta) &= \prod_{i=1}^{n_s} \prod_{j=1}^{n_w} \frac{(\lambda_{ij}(\Theta) n_{ij})^{Y_{ij}} e^{\lambda_{ij}(\Theta) n_{ij}}}{Y_{ij}!} \\ &= \prod_{i=1}^{n_s} \prod_{j=1}^{n_w} \frac{\{\exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon) n_{ij}\}^{Y_{ij}} \exp\{\lambda_{ij}(\Theta) n_{ij}\}}{Y_{ij}!} \\ &\propto \prod_{i=1}^{n_s} \prod_{j=1}^{n_w} \{\exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon) n_{ij}\}^{Y_{ij}} \exp\{\lambda_{ij}(\Theta) n_{ij}\} \\ \log L_Y(\Theta) &= \sum_{i=1}^{n_s} \sum_{j=1}^{n_w} Y_{ij} \log(\lambda_{ij}(\Theta) n_{ij}) + \lambda_{ij}(\Theta) n_{ij} - \log(Y_{ij}!) \\ &\propto \sum_{i=1}^{n_s} \sum_{j=1}^{n_w} Y_{ij} \log(\lambda_{ij}(\Theta) n_{ij}) + \lambda_{ij}(\Theta) n_{ij} \\ &\propto \sum_{i=1}^{n_s} \sum_{j=1}^{n_w} Y_{ij} (\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon + \log n_{ij}) + \exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon) n_{ij} \end{aligned}$$

3.1.4 Posterior

The posterior distribution is proportional to the product of the likelihood and prior:

$$\begin{aligned} g(\Theta | Y) &\propto L_Y(\Theta) \pi(\Theta) \\ &\propto \prod_{i=1}^{n_s} \prod_{j=1}^{n_w} \frac{(\lambda_{ij}(\Theta) n_{ij})^{Y_{ij}} e^{\lambda_{ij}(\Theta) n_{ij}}}{Y_{ij}!} \exp\left\{-\frac{\alpha^2 - \sum_{i=1}^k \beta_k^2 - \sigma_u^2 - \sigma_\epsilon^2}{2 \cdot 10^2}\right\} \frac{1}{\sigma_u \sigma_\epsilon} \exp\left\{-\frac{\sum_{i=1}^{n_s} u_i^2}{2\sigma_u^2}\right\} \exp\left\{-\frac{\epsilon^2}{2\sigma_\epsilon^2}\right\} \\ &\propto \prod_{i=1}^{n_s} \prod_{j=1}^{n_w} \{\exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon) n_{ij}\}^{Y_{ij}} \exp\{\lambda_{ij}(\Theta) n_{ij}\} \\ &\quad \cdot \exp\left\{-\frac{\alpha^2 - \sum_{i=1}^k \beta_k^2 - \sigma_u^2 - \sigma_\epsilon^2}{2 \cdot 10^2}\right\} \frac{1}{\sigma_u \sigma_\epsilon} \exp\left\{-\frac{\sum_{i=1}^{n_s} u_i^2}{2\sigma_u^2}\right\} \exp\left\{-\frac{\epsilon^2}{2\sigma_\epsilon^2}\right\} \end{aligned}$$

$$\begin{aligned}
logg(\Theta|Y) &\propto logL_Y(\Theta) + log\pi(\Theta) \\
&\propto \sum_{i=1}^{n_s} \sum_{j=1}^{n_w} Y_{ij}(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon + logn_{ij}) + exp(\alpha + \beta X_{ij} + \gamma P_{ij} + \delta E_{ij} + u_i + \epsilon)n_{ij} \\
&+ \frac{-\alpha^2 - \sum_{i=1}^k \beta_k^2 - \sigma_u^2 - \sigma_\epsilon^2}{2 \cdot 10^2} - log(\sigma_u) - log(\sigma_\epsilon) - \frac{\sum_{i=1}^{n_s} u_i^2}{2\sigma_u^2} - \frac{\epsilon^2}{2\sigma_\epsilon^2}
\end{aligned}$$

3.2 Component-wise MH Algorithm

The Component-wise Metropolis-Hastings (CMH) algorithm is a variant of the Metropolis-Hastings(MH) algorithm used for sampling from high-dimensional probability distributions. The CMH algorithm updates each component of a high-dimensional parameter vector separately, rather than updating the entire vector jointly.

The main idea behind the CMH algorithm is to treat each component of the parameter vector as a separate one-dimensional distribution, and update it using a one-dimensional MH proposal. At each iteration of the algorithm, a random component is selected and updated while keeping all other components fixed. The acceptance probability is then computed based on the ratio of the target density at the proposed and current values of the component.

The CMH algorithm can be useful in situations where updating the entire vector jointly is computationally expensive or impractical, such as in Bayesian inference for complex models with high-dimensional parameters.

In this study, there is a high-dimensional parameter vector

$$\Theta = \{\alpha, \beta_1, \dots, \beta_4, \gamma, \delta, u_1, \dots, u_{50}, \epsilon, \sigma_u, \sigma_\epsilon\}$$

The CMH algorithm proceeds as follows:

1. Initialize M chains of length T , with each chain starting from a different initial value of Θ . For each iteration $t = 1, 2, \dots, T$ and for each chain $m = 1, 2, \dots, M$, randomly select a component j from $1, 2, \dots, n$.
2. Propose a new value $\Theta_{j,t}$ for the t -th iteration component j of chain m using a one-dimensional Metropolis-Hastings update. That is, draw a proposal $\Theta_{j,t} \sim q(\cdot|\Theta_{j,t}^{(m)})$, where $q(\cdot|\Theta_{j,t}^{(m)})$ is a proposal distribution centered at the current value $\Theta_{j,t}^{(m)}$ of component j :

$$q(\cdot|\Theta_{j,t}^{(m)}) = q(\cdot|\Theta_{j-1,t}^{(m)}) + w_j * 2 * (r_{j,t} - 0.5)$$

where w_j is the window length, and $r_{j,t}$ is a random number follows $Uniform(0, 1)$.

3. Compute the acceptance probability

$$a_{j,t}(\Theta_{j,t}^{(m)}, \Theta_{j,t}) = \min\left\{1, \frac{p(\Theta_{1,t}^{(m)}, \dots, \Theta_{j-1,t}^{(m)}, \Theta_{j,t}, \Theta_{j+1,t}^{(m)}, \dots, \Theta_n^{(m)})}{p(\Theta_{1,t}^{(m)}, \dots, \Theta_{n,t}^{(m)})} \frac{q(\Theta_{j,t}^{(m)}|\Theta_{j,t})}{q(\Theta_{j,t}|\Theta_{j,t}^{(m)})}\right\}$$

. where $p(\cdot)$ is the target density of the parameter vector Θ .

4. Accept the proposed new value Θ_j with probability $a_{j,t}(\Theta_{j,t}^{(m)}, \Theta_{j,t})$, and set $\Theta_{j,t}^{(m+1)} = \Theta_{j,t}^*$ if the proposal is accepted, and $\Theta_{j,t}^{(m+1)} = \Theta_{j,t}^{(m)}$ otherwise.
5. Repeat steps 2-5 until convergence is achieved.

The convergence of the CMH algorithm can be assessed using the same diagnostics as for standard MCMC algorithms, such as examining the trace plots, autocorrelation plots, and the Gelman-Rubin statistic.

3.3 MCMC Chain Convergence Diagnostics

3.3.1 Diagnostic Plots

The trace plot, also called as chain plot, is a graphical representation of the values of a parameter as a function of the iteration number which shows how the parameter value changes over time in the MCMC chain. Ideally, the trace plot will show a sequence of random and independent samples from the posterior distribution, indicating that the chain has converged to the target distribution and is mixing well. However, if the trace plot shows patterns or trends, such as long runs of increasing or decreasing values or oscillations around a certain value, indicating that the chain has not yet converged or is not mixing well. It also enables to detect potential problems in the sampling process, such as inappropriate step size, or a high correlation, to help improve the efficiency and accuracy of the sampling process.

In an MCMC chain, the histogram is also a graphical representation of the distribution of samples drawn from the posterior distribution. which provides a visual representation of the posterior distribution of the parameter and can also be used to diagnose convergence and mixing issues in the MCMC chain. In an ideal scenario, the histogram will show a smooth, uni-modal and bell-shaped distribution without visible outliers. However, if the histogram shows multiple modes or significant skewness, it may indicate that the chain has not yet converged or is not mixing well.

The autocorrelation plot is another useful diagnostic plot which shows how the values of the parameter at one point in the chain are correlated with consecutive samples in this chain. An ideal autocorrelation function will rapidly decrease to zero as the lag increases, indicating that the parameter values are uncorrelated and the chain has converged to the target distribution. If the autocorrelation function remains high for several lags, it suggests that the chain has not mixed well, and the samples are highly correlated. In this case, additional adjustments should be conducted.

3.3.2 Gelman-Rubin Statistic

The Gelman-Rubin statistic, also known as R-hat, is a convergence diagnostic statistic widely used in MCMC simulations. It helps to assess the convergence of multiple chains run in parallel by comparing the within-chain and between-chain variances and conducting a family of tests.

The calculation of statistic is defined as[1]:

Let $x_1^{(j)}, x_2^{(j)}, \dots$ be samples from the j th Markov chain, and suppose there are J chains run in parallel with different starting values.

1. For each chain, first discard D values as "burn-in" and keep the remaining L values, $x^{(j)}D, x^{(j)}D + 1, \dots, x_{D+L-1}^{(j)}$. For example, you might set $D = L$.
2. Formulas

Chain mean:

$$\bar{x}_j = \frac{1}{L} \sum t = 1^L x_t^{(j)}$$

Grand mean:

$$\bar{x}_\cdot = \frac{1}{J} \sum j = 1^J \bar{x}_j$$

Between chain variance:

$$B = \frac{L}{J-1} \sum_{j=1}^J (\bar{x}_j - \bar{x}_\cdot)^2$$

Within-chain variance:

$$s_j^2 = \frac{1}{L-1} \sum t = 1^L (x_t^{(j)} - \bar{x}_j)^2$$

and

$$W = \frac{1}{J} \sum_{j=1}^J s_j^2$$

3. The Gelman-Rubin statistic is then

$$R = \frac{\frac{L-1}{L}W + \frac{1}{L}B}{W}$$

We can see that as $L \rightarrow \infty$ and as $B \rightarrow 0$, R approaches the value of 1. One can then reason that we should run our chains until the value of R is close to 1, say in the neighborhood of 1.1[2].

The Gelman-Rubin statistic is a ratio and hence unit-free, making it a simple summary for any MCMC sampler. In addition, it can be implemented without first specifying a parameter that is to be estimated, unlike Monte Carlo standard errors. Therefore, it can be a useful tool for monitoring a chain before any specific decisions about what kinds of inferences will be made from the model.

4 Result

4.1 Convergence of MCMC Chain

Figure 2-7 show the trace plots and histograms for all parameters using the final window length after adjustment shown in Table 2, indicating good convergence of some parameters, such as β_1 , β_2 , β_3 , β_4 , γ , σ_s and all the us , which means the samples are likely to be representative of our target distribution and the MCMC chain mixes well.

Table 2: Final window length after adjustment

α	β_1	β_2	β_3	β_4	γ	δ	$u_{1-u_{50}}$	ϵ	σ_u	σ_e
0.1	0.1	0.1	0.1	0.1	0.1	0.1	1	0.1	0.1	10

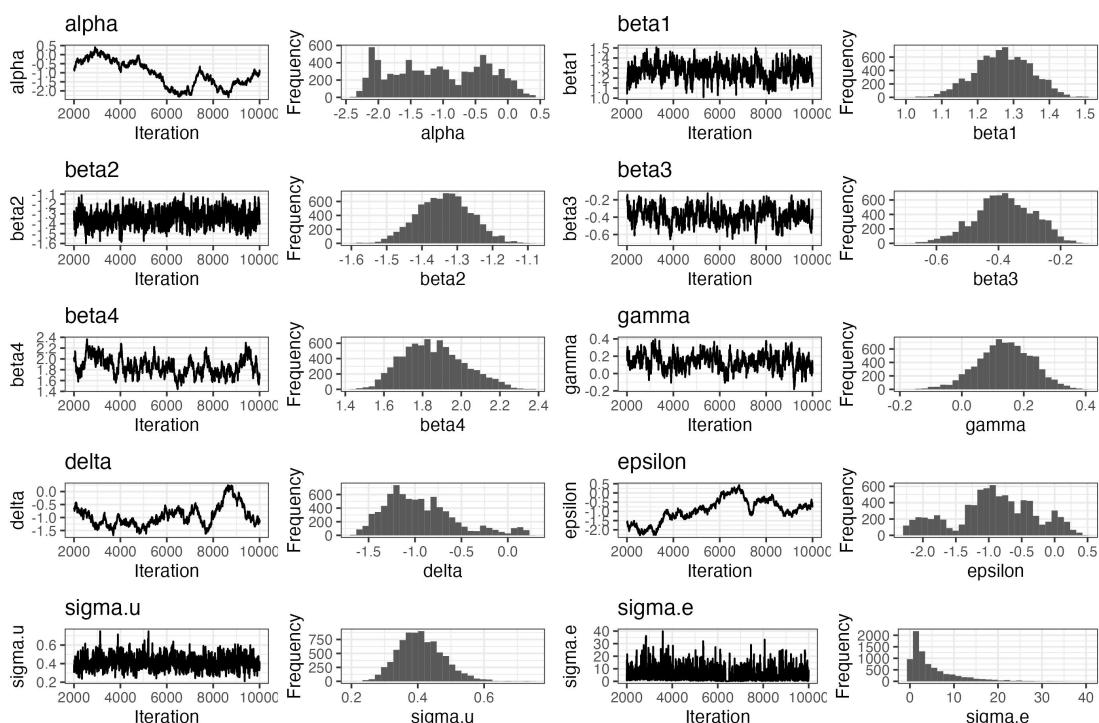


Figure 2: Trace plot for α , β s, γ , δ , ϵ and σ s

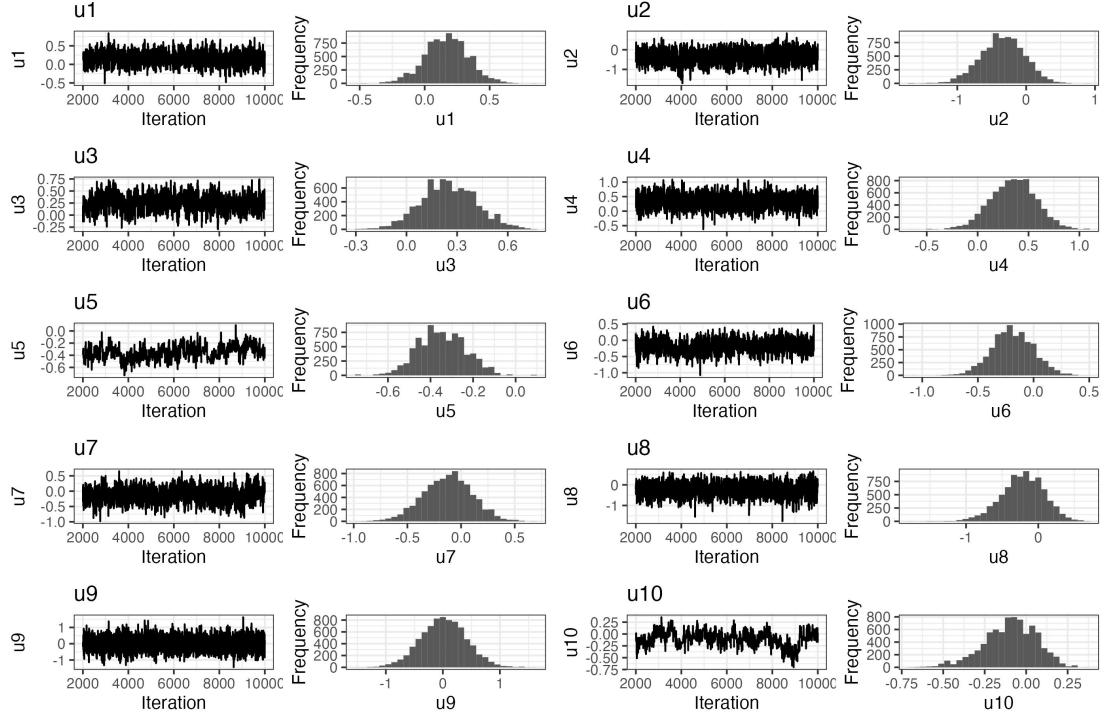


Figure 3: Trace plot for u_1-u_{10}

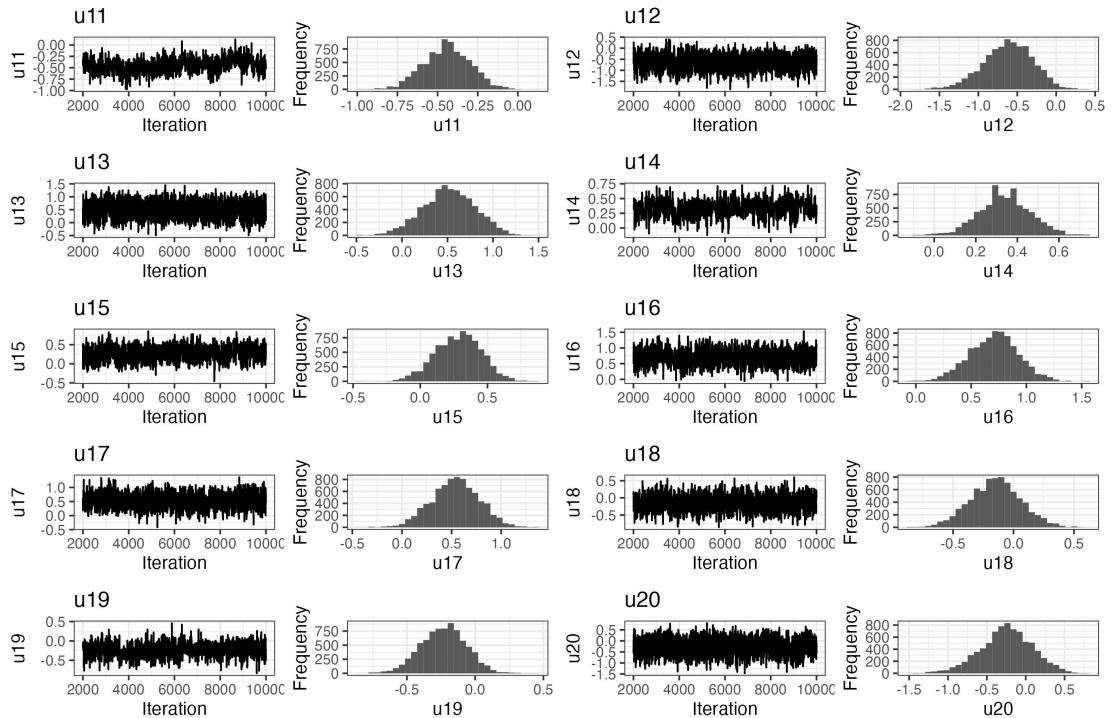


Figure 4: Trace plot for $u_{11}-u_{20}$

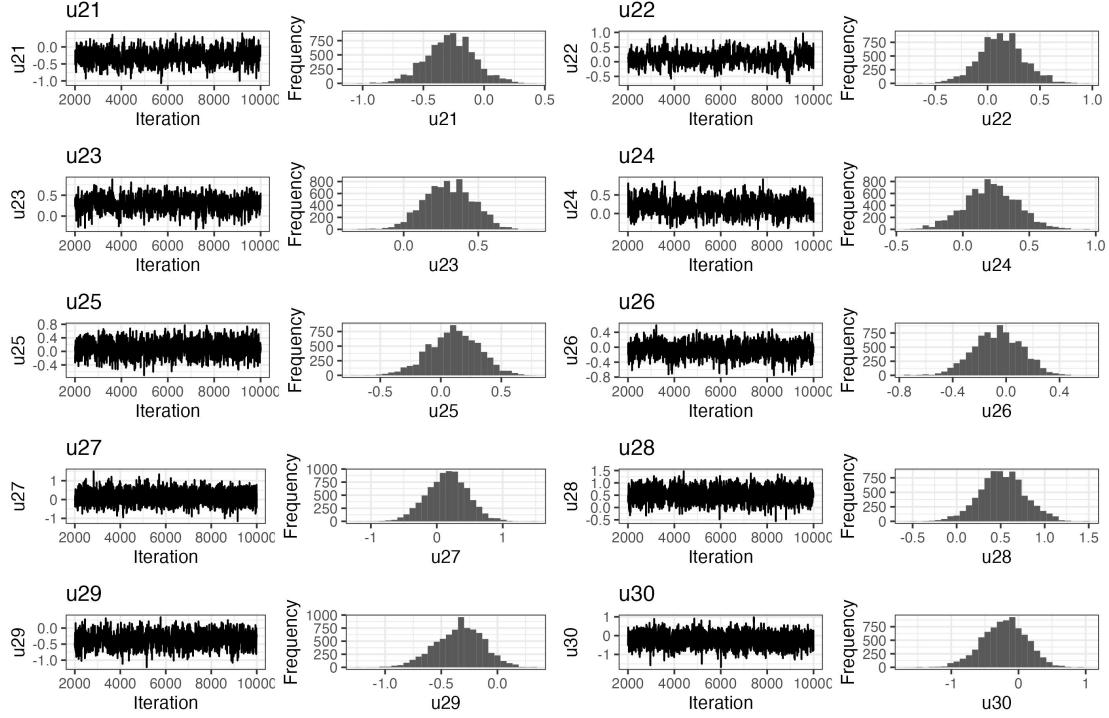


Figure 5: Trace plot for $u_{21}-u_{30}$

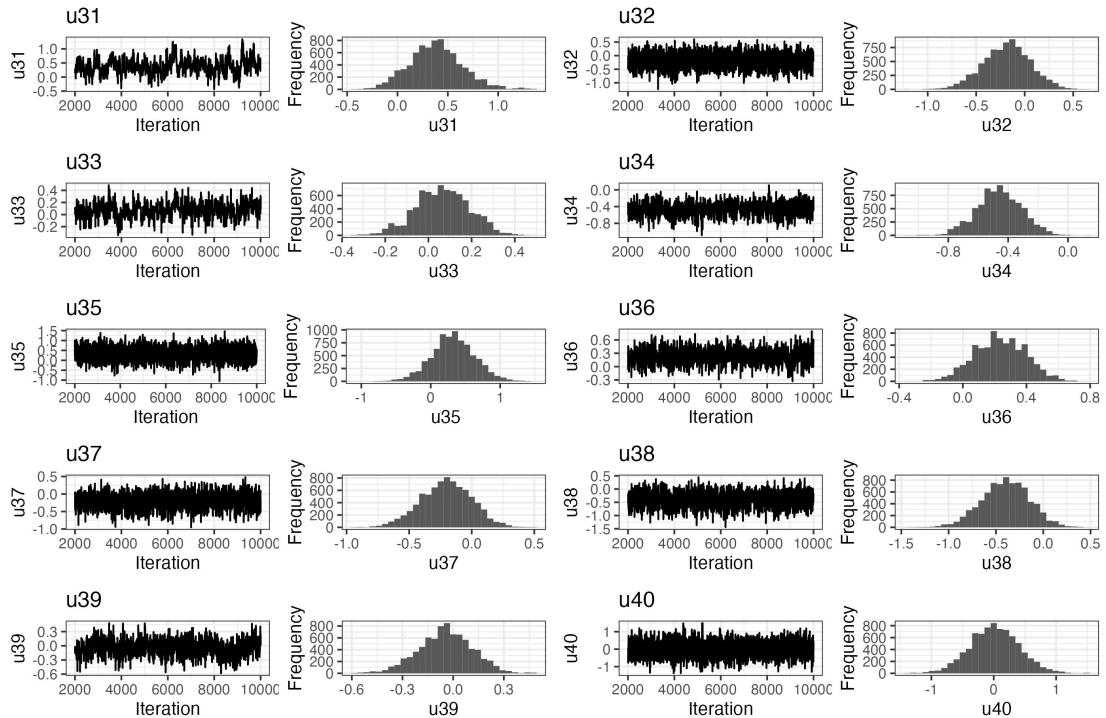


Figure 6: Trace plot for $u_{31}-u_{40}$

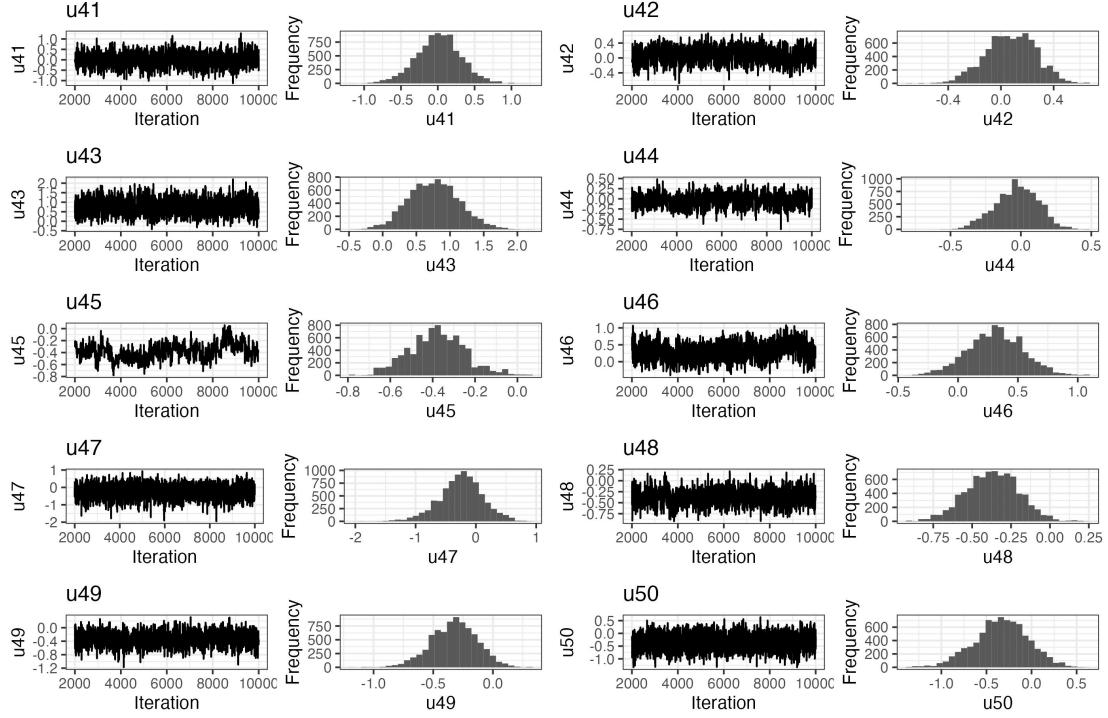


Figure 7: Trace plot for $u_{41}-u_{50}$



Figure 8: Autocorrelation plot for all parameters

Figure 8 is the autocorrelation plot for each parameter. In this plot, most parameters

have a rapid decrease as the lag increases, such as all the u s, which indicates that the samples are independent to the consecutive ones and the Markov chain has mixed well. However, the autocorrelation of some parameters remains horizontal, which suggests that the samples are highly correlated and the chain has not mixed well.

Three chains are generated in parallel using three sets of starting values shown in Table 3. Figure 9-13 shows trace plots for each parameter merged together from three chains. Most of the parameters converge after about 1000 iterations. For these parameters that are able to converge in the end, the selection of starting value does not really influence the final convergence result.

Table 3: Starting values

α	β_1	β_2	β_3	β_4	γ	δ	$u_{1-u_{50}}$	ϵ	σ_u	σ_e
0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1
1	1	1	1	1	1	1	1	1	1	1
-1	-1	-1	-1	-1	-1	-1	-1	-1	0.1	0.1

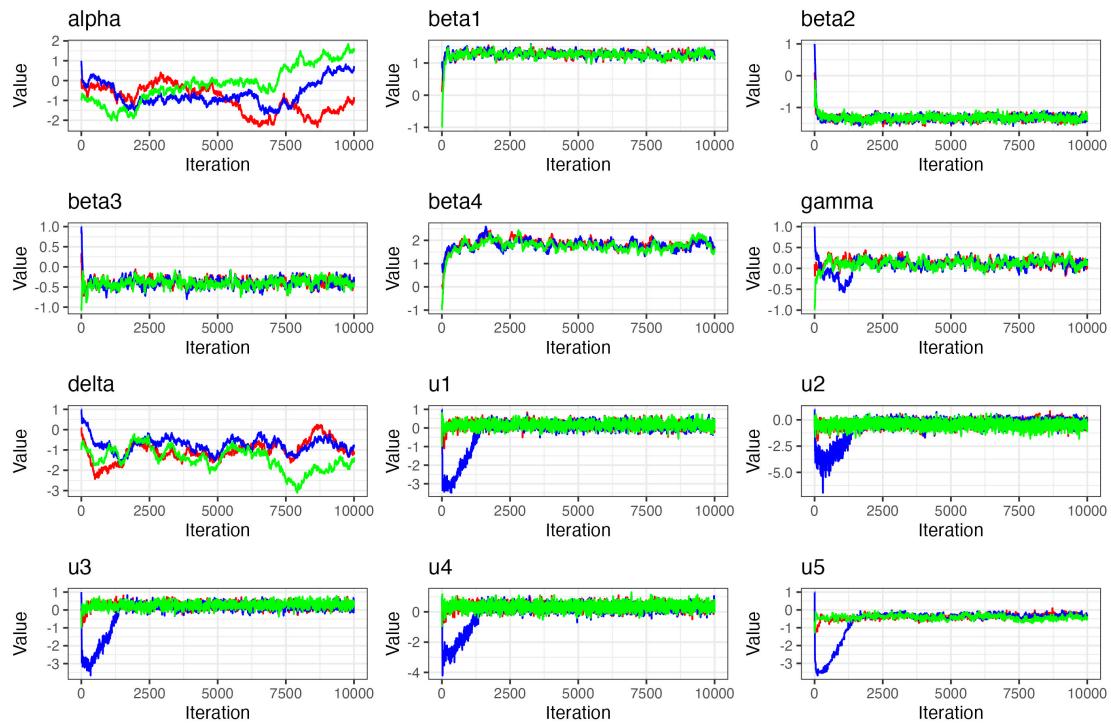


Figure 9: Merged trace plot for α , β s, γ , δ and u_1-u_5

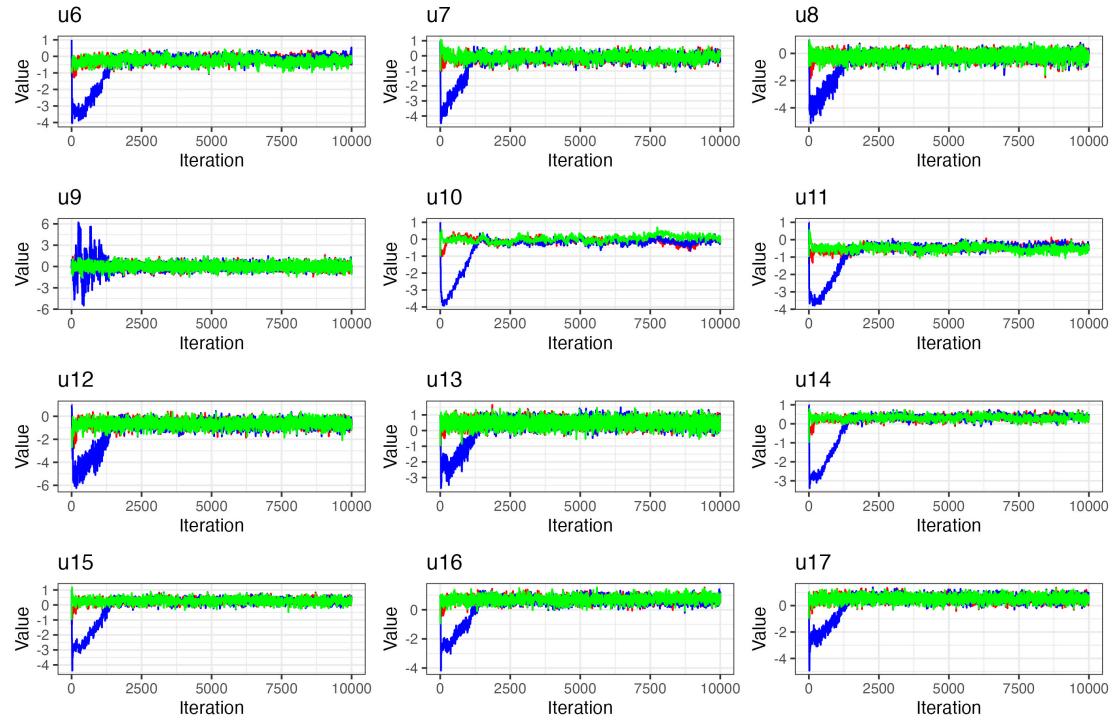


Figure 10: Merged trace plot for u_6-u_{17}

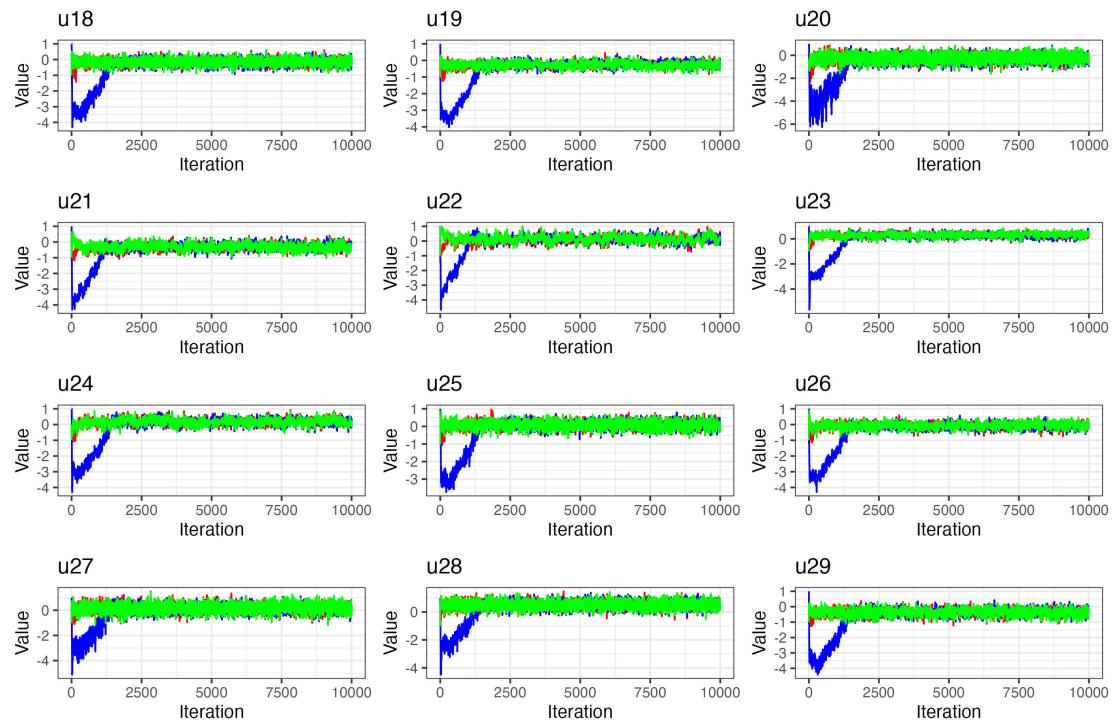


Figure 11: Merged trace plot for $u_{18}-u_{29}$

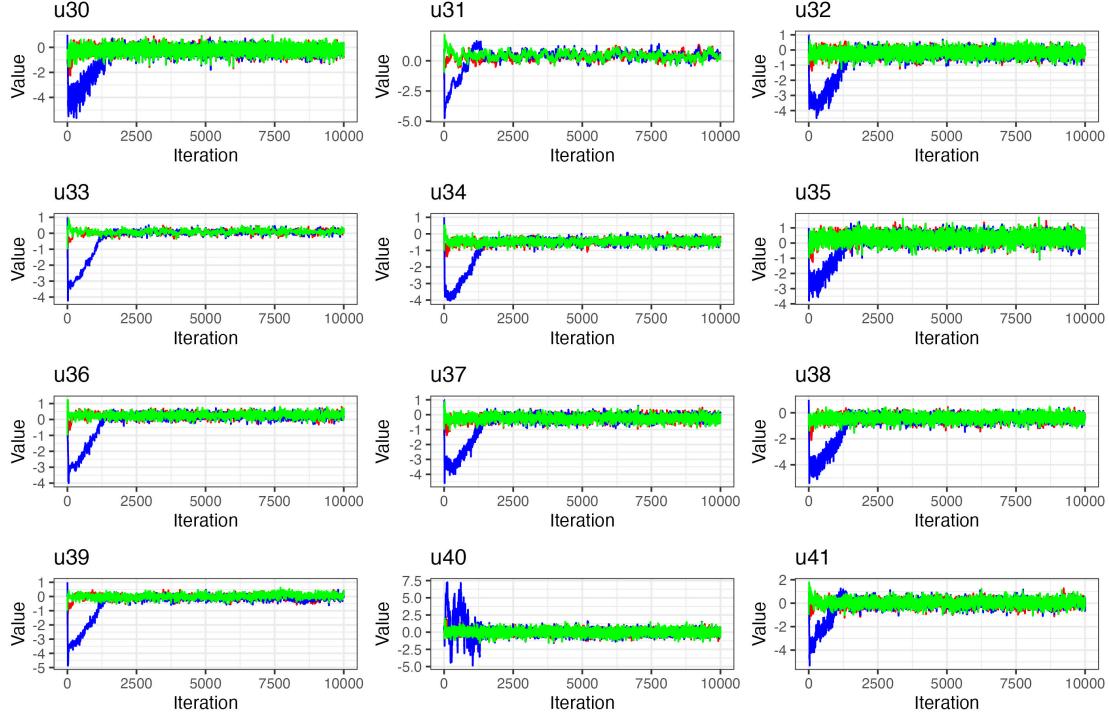


Figure 12: Merged trace plot for u_{30} - u_{41}

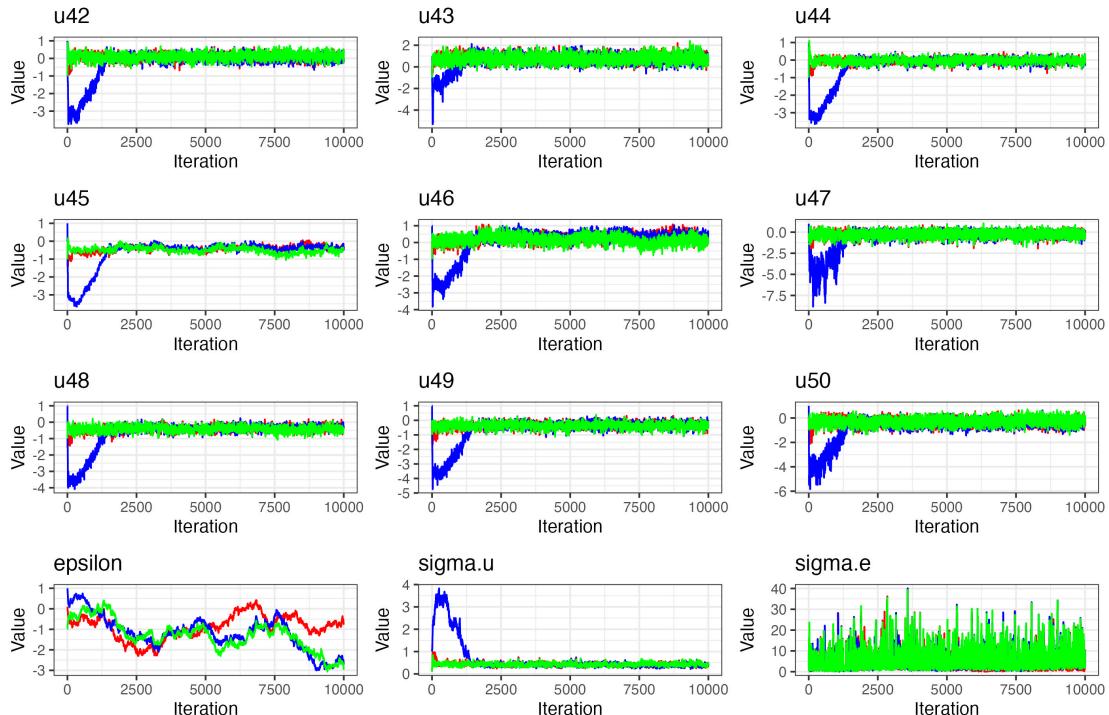


Figure 13: Merged trace plot for u_{42} - u_{50} , ϵ and σ_u

Table 4 shows the Gelman-Rubin statistics calculated based on definitions. Since a statistic below 1.1 indicates convergence of MCMC chain, 57 over 60 parameters converge and only 3 of them do not converge, including α , δ and u_{10} .

4.2 Posterior Distribution

Table 5 gives the posterior summaries for all the parameters of interest. It shows that mobility changes in parks and transit stations and state elderly percentage are negatively associated with the infection rate, while mobility changes in retail and recreation places, government response and state population density are positively associated with the infection rate. The absolute mean value of β_1 , β_2 and β_4 are greater than 1, suggesting that mobility change in retail, recreation places and parks, and the government response index have relatively large effect size compared to the rest.

Table 4: Gelman-R Rubin Statistics

Para	Stat	Converge	Para	Stat	Converge	Para	Stat	Converge
α	1.20	FALSE	u_{14}	1.06	TRUE	u_{34}	1.08	TRUE
β_1	1.01	TRUE	u_{15}	1.07	TRUE	u_{35}	1.05	TRUE
β_2	1.00	TRUE	u_{16}	1.08	TRUE	u_{36}	1.08	TRUE
β_3	1.00	TRUE	u_{17}	1.06	TRUE	u_{37}	1.07	TRUE
β_4	1.02	TRUE	u_{18}	1.07	TRUE	u_{38}	1.08	TRUE
γ	1.04	TRUE	u_{19}	1.07	TRUE	u_{39}	1.10	TRUE
δ	1.36	FALSE	u_{20}	1.07	TRUE	u_{40}	1.01	TRUE
u_1	1.09	TRUE	u_{21}	1.04	TRUE	u_{41}	1.04	TRUE
u_2	1.04	TRUE	u_{22}	1.04	TRUE	u_{42}	1.09	TRUE
u_3	1.10	TRUE	u_{23}	1.09	TRUE	u_{43}	1.04	TRUE
u_4	1.08	TRUE	u_{24}	1.07	TRUE	u_{44}	1.08	TRUE
u_5	1.05	TRUE	u_{25}	1.07	TRUE	u_{45}	1.04	TRUE
u_6	1.05	TRUE	u_{26}	1.09	TRUE	u_{46}	1.03	TRUE
u_7	1.05	TRUE	u_{27}	1.07	TRUE	u_{47}	1.06	TRUE
u_8	1.07	TRUE	u_{28}	1.06	TRUE	u_{48}	1.06	TRUE
u_9	1.00	TRUE	u_{29}	1.07	TRUE	u_{49}	1.07	TRUE
u_{10}	1.14	FALSE	u_{30}	1.06	TRUE	u_{50}	1.08	TRUE
u_{11}	1.05	TRUE	u_{31}	1.02	TRUE	ϵ	1.05	TRUE
u_{12}	1.07	TRUE	u_{32}	1.08	TRUE	σ_u	1.07	TRUE
u_{13}	1.06	TRUE	u_{33}	1.07	TRUE	σ_e	1.00	TRUE

Table 5: Posterior summaries

Statistic	N	Mean	St. Dev.	Min	Max
β_1	5,001	1.264	0.077	1.008	1.500
β_2	5,001	-1.330	0.076	-1.576	-1.090
β_3	5,001	-0.386	0.095	-0.702	-0.124
β_4	5,001	1.823	0.149	1.437	2.301
γ	5,001	0.136	0.089	-0.182	0.377
δ	5,001	-0.824	0.414	-1.631	0.246

β_1 : change in mobility trend in retail and recreation places; β_2 : change in mobility trend in parks; β_3 : change in mobility trend in transit stations; β_4 : government response index; γ : coefficient for state population density; δ : coefficient for state elderly percentage.

4.3 Credible Intervals

Bayesian posterior uncertainty intervals, often referred to as credible intervals, are used to approximate the variability of parameters from MCMC draws.

Figure 15 and Table 6 show that the intervals are narrow for β_1 , β_2 , β_3 , β_4 , γ , but wide for δ . It means that the effect of government interventions, mobility changes and population density are more reliable compared with that of the elderly percentage on the infection rate. The increase in the weekly average percentage change in mobility trends for retail and recreation places will increase the number of new infections, while more changes in mobility trends for parks and transit stations tend to lower the infection rate. The credible interval of β_4 and γ contain 0, suggesting that the effects of population density and elderly percentage are not significant. Therefore, we should focus more on government interventions and mobility changes when giving recommendations for public health policies.

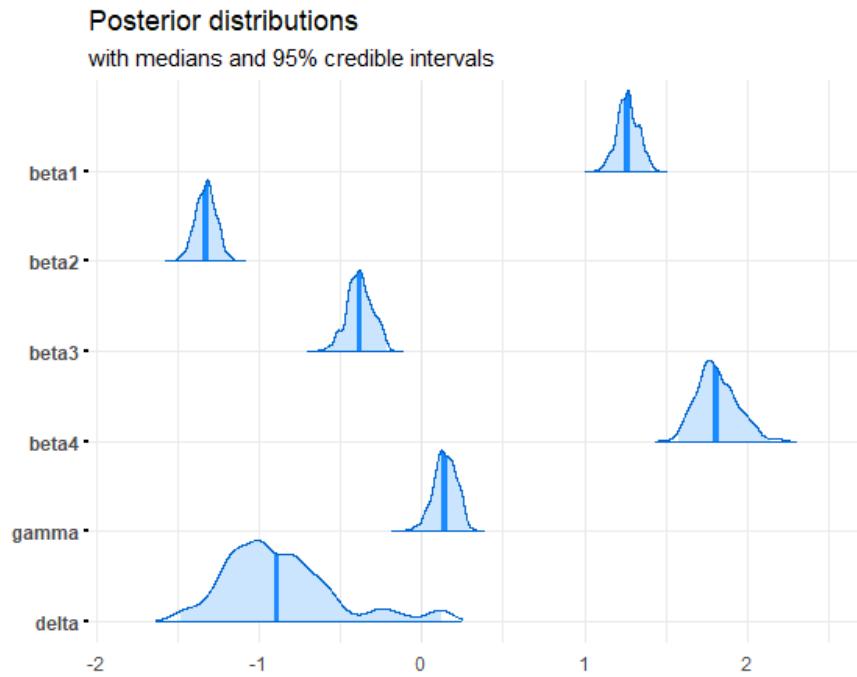


Figure 14: Posterior distributions of parameters of interest

Table 6: 95% credible intervals

	2.5%	97.5%
β_1	1.111	1.411
β_2	-1.478	-1.178
β_3	-0.579	-0.217
β_4	1.575	2.146
γ	-0.053	0.296
δ	-1.478	0.123

5 Discussion

5.1 Interpretation

In terms of model building, we found that collinearity weakens MCMC parameter convergence. This stems from the observation that the model convergence was poor when all the highly correlated indices were included, but it improved significantly when only government response index was kept. We also found that scaling the predictors helps the MCMC chains to converge. In summary, data inspection and processing steps are essential for model fitting.

In terms of results interpretation, we first note that some results were counterintuitive at the first sight. However, some potential explanations can be provided through the temporal relationships between the pandemic and the covariates.

For example, recall that we found some variational effects of changes in mobility trends in COVID infection cases, that is, we found a positive association between cases and mobility changes in retail, but a negative association between cases and mobility changes in transit and parks. This can be explained by looking at the trends of these mobility changes and COVID cases over weeks. First, note that the 2020 COVID pandemic has three waves based on the figure: the first two were relatively small in a few states, approximately from week 8 to week 11 and from week 20 to week 26. The third wave started around week 36, affecting most states. Corresponding this trend to mobility changes in parks and transit stations, we can see that during the third wave, both parks and transit mobility decreased a lot in magnitude ($\sim 200\%$ for parks; $\sim 100\%$ for transit), which may have resulted in the negative association between COVID cases and parks and transit mobility. However, retail mobility increased drastically during the first and second wave, and only decreased in a much smaller magnitude ($\sim 20\%$) than parks and transit mobility during the third wave of the outbreak. A more contextual explanation for this negative association is that there were decreasing levels of outdoor activities and traveling and an increasing number of remote work (less commute) when there was an outbreak. In comparison, retail activities like grocery shopping are essential for most, and some even felt an increasing need to store more necessities. In short, the pandemic has different impacts on essential vs. less essential activities, resulting in a differential association between COVID cases and mobility changes.

The positive association between infection and government response index can be explained similarly - that is, the worse the outbreak, the higher level of government response, as indicated by the overall increase in the government response index as the pandemic progressed.

As for the non-significant association between COVID cases and state-level population density and percent of elderly population, it can be confounded by other factors such as the economy of the state. It is possible that after adjusting for relevant confounders, the relationship between COVID cases and state demographics can be clearer and more informative for making public health decisions.

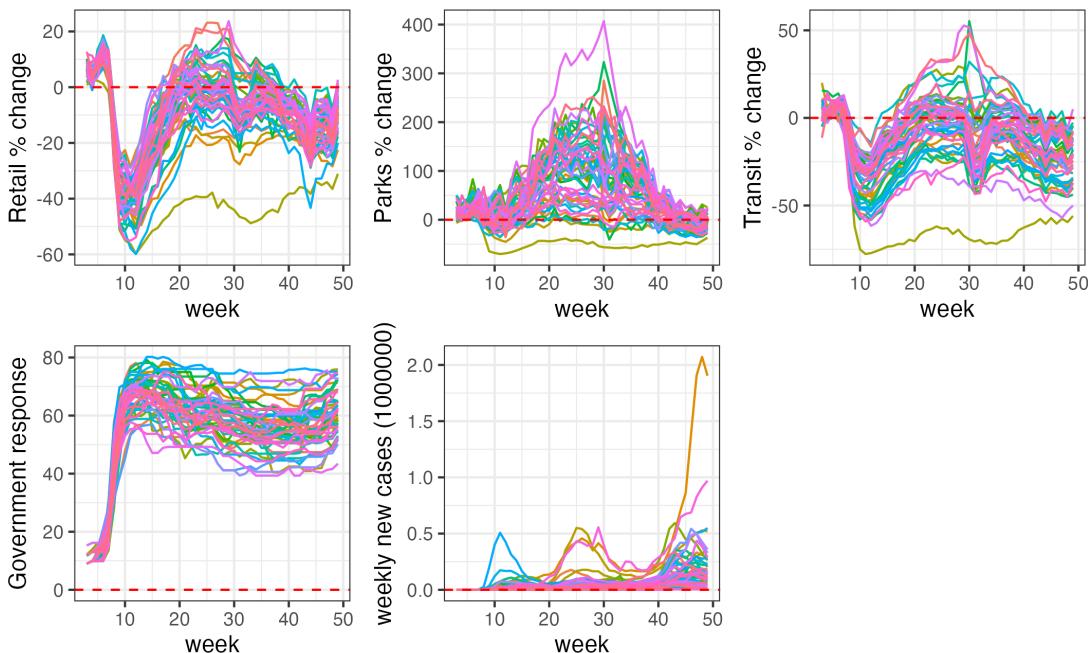


Figure 15: Trends of COVID cases, retail % change from baseline, parks % change from baseline, transit % change from baseline, and government response index over week

5.2 Public Health Policy Recommendations

Based on the results, we propose several public health policy recommendations for COVID containment. First, it is important to restrict unnecessary mobility at early stages of the outbreak, because waiting until the outbreak has spread may render restrictions less effective, as we have witnessed in 2020. It is also important to improve accessibility to COVID tests, especially in less economically developed states. This is because the recorded number of infections in each state might vary in terms of accuracy because of differential access and distribution of COVID testing. Another recommendation is to research more on confounders between infected cases and relevant indices, as well as state demographics. This helps develop a better statistical model that quantifies the effects of more factors on COVID infection.

5.3 Suggestions on Further Research

The unexpected negative correlation between certain parameter estimates and weekly new case counts suggests that there may be factors at play that are not fully understood or accounted for in the analysis. To investigate the causal relationships between the disease and the intervention, it may be necessary to explore additional variables that could be affecting the outcome. For example, the time lag in government response may be of interest. Also, it is important to consider the interactions among different factors to identify potential risks and effectiveness of interventions. For example, The mobility and government response could be investigated on how they affect each other and influence the overall outcome. In all, further research is needed for better public health outcomes, and helps to advance our knowledge of disease dynamics and intervention strategies.

6 Conclusion

In summary, in our project, we developed a bayesian hierarchical poisson regression model by writing down the log of the prior, likelihood, and posterior density for 60 parameters. We monitored the convergence of MCMC chains through diagnostic plots to ensure adequate model fitting. By constructing 95% credible intervals for all parameters, we were able to provide interpretations and potential explanations for these parameters and think of public health policy recommendations.

Acknowledgement

All the members contributed equally to the project.

References

- [1] Stephen P. Brooks and Andrew Gelman. “General Methods for Monitoring Convergence of Iterative Simulations”. In: *Journal of Computational and Graphical Statistics* 7.4 (1998), pp. 434–455.
- [2] Andrew Gelman and Donald B. Rubin. “Inference from Iterative Simulation Using Multiple Sequences”. In: *Statistical Science* 7.4 (1992), pp. 457–472.

Code Implementation of MCMC

```
# R implementation of MH algorithm:  
# log likelihood  
LikelihoodFunction <- function(param){  
    alpha <- param[1]  
    beta1 <- param[2]  
    beta2 <- param[3]  
    beta3 <- param[4]  
    beta4 <- param[5]  
    gamma <- param[6]  
    delta <- param[7]  
    u <- param[8:(7+n.state)]  
    state.u <- u[state]  
    epsilon <- param[8+n.state]  
    lambda <- exp(alpha + beta1 * X1 + beta2 * X2 + beta3 * X3 + beta4 * X4  
    + gamma * P + delta * E + state.u + epsilon)  
    loglikelihoods <- sum(y * log(lambda*n) - lambda*n, na.rm = TRUE)  
    return(loglikelihoods)  
}  
  
# log prior  
LogPriorFunction <- function(param){  
    alpha <- param[1]  
    beta1 <- param[2]  
    beta2 <- param[3]  
    beta3 <- param[4]  
    beta4 <- param[5]  
    gamma <- param[6]  
    delta <- param[7]  
    u <- param[8:(7+n.state)]  
    state.u <- u[state]  
    epsilon <- param[8+n.state]  
    s.u <- param[9+n.state]  
    s.epsilon <- param[10+n.state]  
    if (min(s.u, s.epsilon) <= 0){  
        return(-Inf)  
    }  
    else{  
        logprior = (-alpha^2-beta1^2-beta2^2-beta3^2-beta4^2-s.u^2-s.epsilon^2)/200  
        -sum(u^2)/(2*s.u^2)-epsilon^2/(2*s.epsilon^2)-n.state*log(s.u)-log(s.epsilon)  
        return(logprior)  
    }  
}  
  
# log posterior  
PosteriorFunction <- function(param){  
    return (LikelihoodFunction(param) + LogPriorFunction(param))  
}  
  
# component-wise MH algorithm  
RunMetropolisMCMC <- function(startvalue, iterations, postfun, a){
```

```

n.parm = length(startvalue)
chain <- array(dim=c(iterations + 1, n.parm))
chain[1, ] <- startvalue
for (i in 1:iterations){
  curr <- chain[i,]
  set.seed(i)
  print(i)
  for (p in 1:n.parm){
    prop <- curr
    prop[p] <- curr[p] + 2 * (runif(1) - 0.5) * a[p]
    if(log(runif(1)) < postfun(prop) - postfun(curr)){
      curr[p] <- prop[p]
    }
  }
  chain[i+1, ] <- curr
}
return(chain)
}

# Gelman-Rubin statistic
gelman.rubin.stat <- function(chains){
  pure.chain <- chains[,-1]
  L <- nrow(pure.chain)
  J <- ncol(pure.chain)
  # chain mean
  chain.mean <- apply(pure.chain, 2, mean)
  # grand mean
  grand.mean <- mean(chain.mean)
  # b/w chain variance
  B <- L / (J-1) * sum((chain.mean-grand.mean)^2)
  # within chain variance
  s2_j <- 1 / (J-1) * apply(pure.chain, 2, var)

  W <- sum(s2_j)
  # Gelman-Rubin statistic
  R <- (L/(L-1)*W + 1/L*B) / W
  R
}

}

```